

Laboratorio 02 – Series de Tiempo LSTM

- **CC3084 – Data Science**, Semestre II 2025
- **Integrantes:** Diego Valenzuela 22309, Gerson Ramirez 22281
- **Guatemala, Julio 2025**

Repositorio:

<https://github.com/Diegoval-Dev/DC-Lab2>

Codebook Docs:

https://docs.google.com/document/d/19wPL80Dss0d-sQldEt3_nUyNCderHo_rLaK94Q4Nuuc/edit?usp=sharing

```
In [118... #%pip install torch matplotlib pandas sklearn
```

```
In [119... import pandas as pd
import numpy as np
import torch
import torch.nn as nn
from torch.utils.data import TensorDataset, DataLoader
from sklearn.preprocessing import MinMaxScaler
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
from torch.utils.data import TensorDataset, DataLoader

# DATA GENERAL PARA CUALQUIERA
df_imp = pd.read_csv('./data/importacion.csv', parse_dates=['fecha'],
df_cons = pd.read_csv('./data/consumo.csv', parse_dates=['fecha'],

window_size = 3
```

```
In [120... def create_dataset(series, window_size=12):
    """
    series: array unidimensional de valores (numpy array)
    window_size: número de pasos de tiempo para cada muestra X
    Devuelve:
        X: array de forma (n_samples, window_size, 1)
        y: array de forma (n_samples, 1)
    """
    X, y = [], []
    for i in range(len(series) - window_size):
        X.append(series[i : i + window_size])
        y.append(series[i + window_size])
```

```
X = np.array(X).reshape(-1, window_size, 1)
y = np.array(y).reshape(-1, 1)
return X, y
```

Evaluación de Series de Tiempo

Gasolina Regular Importaciones


```
In [121... ts_ir = df_imp.set_index('fecha')['gasolina regular'] \
            .resample('M').mean().dropna()

print(f"Regular_imp - Inicio: {ts_ir.index.min().date()}, Fin: {ts_ir.

plt.figure(figsize=(10,3))
plt.plot(ts_ir, color='tab:blue')
plt.title('Gasolina Regular - Importación')
plt.show()

values_ir = ts_ir.values
```

Regular_imp - Inicio: 2001-01-31, Fin: 2024-01-31, Frecuencia: Mensual
/var/folders/kd/4y1c0b3j1273_pm_gb0hfhdw0000gn/T/ipykernel_92772/288097
8788.py:2: FutureWarning: 'M' is deprecated and will be removed in a fu
ture version, please use 'ME' instead.
.resample('M').mean().dropna()



Gasolina Regular Consumos


```
In [122... ts_cr = df_cons.set_index('fecha')['gasolina regular'] \
            .resample('M').mean().dropna()

print(f"Regular_cons - Inicio: {ts_cr.index.min().date()}, Fin: {ts_cr

plt.figure(figsize=(10,3))
plt.plot(ts_cr, color='tab:green')
plt.title('Gasolina Regular - Consumo')
plt.show()
```

```
values_cr = ts_cr.values
```

Regular_cons – Inicio: 2000-01-31, Fin: 2024-01-31, Frecuencia: Mensual
/var/folders/kd/4y1c0b3j1273_pm_gb0hfhdw0000gn/T/ipykernel_92772/375153888.py:2: FutureWarning: 'M' is deprecated and will be removed in a future version, please use 'ME' instead.
.resample('M').mean().dropna()



Gasolina Superior Importaciones


```
In [123... ts_cr = df_imp.set_index('fecha')['gasolina superior'] \
            .resample('M').mean().dropna()

print(f"Superior_imp – Inicio: {ts_cr.index.min().date()}, Fin: {ts_cr.index.max().date()}")

plt.figure(figsize=(10,3))
plt.plot(ts_cr, color='tab:green')
plt.title('Gasolina Superior – Importación')
plt.show()

values_cr = ts_cr.values
```

Superior_imp – Inicio: 2001-01-31, Fin: 2024-01-31, Frecuencia: Mensual
/var/folders/kd/4y1c0b3j1273_pm_gb0hfhdw0000gn/T/ipykernel_92772/1285877573.py:2: FutureWarning: 'M' is deprecated and will be removed in a future version, please use 'ME' instead.
.resample('M').mean().dropna()



Diesel Importacion

```
In [124... ts_cr = df_imp.set_index('fecha')['diesel'] \
            .resample('M').mean().dropna()

print(f"Diesel_imp - Inicio: {ts_cr.index.min().date()}, Fin: {ts_cr.i


plt.figure(figsize=(10,3))
plt.plot(ts_cr, color='tab:green')
plt.title('Diesel - Importación')
plt.show()

values_cr = ts_cr.values
```

/var/folders/kd/4y1c0b3j1273_pm_gb0hfhdw0000gn/T/ipykernel_92772/2453093283.py:2: FutureWarning: 'M' is deprecated and will be removed in a future version, please use 'ME' instead.

```
.resample('M').mean().dropna()
```

Diesel_imp - Inicio: 2001-01-31, Fin: 2024-01-31, Frecuencia: Mensual



Gasolina Superior Consumo

```
In [125... ts_cr = df_cons.set_index('fecha')['gasolina superior'] \
            .resample('M').mean().dropna()

print(f"Superior_cons - Inicio: {ts_cr.index.min().date()}, Fin: {ts_c


plt.figure(figsize=(10,3))
plt.plot(ts_cr, color='tab:green')
plt.title('Gasolina Superior - Consumo')
plt.show()

values_cr = ts_cr.values
```

/var/folders/kd/4y1c0b3j1273_pm_gb0hfhdw0000gn/T/ipykernel_92772/1101280297.py:2: FutureWarning: 'M' is deprecated and will be removed in a future version, please use 'ME' instead.

```
.resample('M').mean().dropna()
```

Superior_cons – Inicio: 2000-01-31, Fin: 2024-01-31, Frecuencia: Mensual



Diesel Consumo

```
In [126... ts_cr = df_cons.set_index('fecha')['diesel bajo azufre'] \
            .resample('M').mean().dropna()

print(f"Diesel_cons – Inicio: {ts_cr.index.min().date()}, Fin: {ts_cr.index.max().date()}")


plt.figure(figsize=(10,3))
plt.plot(ts_cr, color='tab:green')
plt.title('Diesel – Consumo')
plt.show()

values_cr = ts_cr.values
```

/var/folders/kd/4y1c0b3j1273_pm_gb0hfhdw0000gn/T/ipykernel_92772/3015851994.py:2: FutureWarning: 'M' is deprecated and will be removed in a future version, please use 'ME' instead.

```
.resample('M').mean().dropna()
```

Diesel_cons – Inicio: 2000-01-31, Fin: 2024-01-31, Frecuencia: Mensual



Preparación de Datos - Gasolina Regular Importaciones

```
In [127... # Crear dataset para Gasolina Regular Importaciones
X_ir, y_ir = create_dataset(values_ir, window_size)
```

```

print(f"IR → X_ir: {X_ir.shape}, y_ir: {y_ir.shape}")

# División train/test para IR
train_end = '2023-12-31'
test_start = '2024-01-31'

train_ts_ir = ts_ir[:train_end]
test_ts_ir = ts_ir[test_start:]

print(f"Train IR tiene {len(train_ts_ir)} puntos, Test IR tiene {len(test_ts_ir)} puntos")
print(f"IR train rango: {train_ts_ir.index[0].date()} – {train_ts_ir.index[-1].date()}")

# Recrear dataset con datos de entrenamiento únicamente
X_ir, y_ir = create_dataset(train_ts_ir.values, window_size)
print(f"Después de split y ventanas → IR: X={X_ir.shape}, y={y_ir.shape}")

```

IR → X_ir: (21, 3, 1), y_ir: (21, 1)
 Train IR tiene 23 puntos, Test IR tiene 1
 IR train rango: 2001-01-31 – 2023-01-31
 Después de split y ventanas → IR: X=(20, 3, 1), y=(20, 1)

Preparación de Datos - Gasolina Regular Consumos

In [128...

```

# Crear dataset para Gasolina Regular Consumos
X_cr, y_cr = create_dataset(values_cr, window_size)
print(f"CR → X_cr: {X_cr.shape}, y_cr: {y_cr.shape}")

# División train/test para CR
train_ts_cr = ts_cr[:train_end]
test_ts_cr = ts_cr[test_start:]

print(f"Train CR tiene {len(train_ts_cr)} puntos, Test CR tiene {len(test_ts_cr)} puntos")
print(f"CR train rango: {train_ts_cr.index[0].date()} – {train_ts_cr.index[-1].date()}")

# Recrear dataset con datos de entrenamiento únicamente
X_cr, y_cr = create_dataset(train_ts_cr.values, window_size)
print(f"Después de split y ventanas → CR: X={X_cr.shape}, y={y_cr.shape}")

```

CR → X_cr: (22, 3, 1), y_cr: (22, 1)
 Train CR tiene 24 puntos, Test CR tiene 1
 CR train rango: 2000-01-31 – 2023-01-31
 Después de split y ventanas → CR: X=(21, 3, 1), y=(21, 1)

Entrenamiento LSTM - Gasolina Regular Importaciones

In [129...

```

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Usando dispositivo: {device}")

# Preparar datos para PyTorch - IR

```

```

tensor_X_ir = torch.tensor(X_ir, dtype=torch.float32).to(device)
tensor_y_ir = torch.tensor(y_ir, dtype=torch.float32).to(device)

batch_size = 16
loader_ir = DataLoader(TensorDataset(tensor_X_ir, tensor_y_ir),
                        batch_size=batch_size, shuffle=True)

# Definir modelo LSTM
class LSTMModel(nn.Module):
    def __init__(self, hidden_size, dropout):
        super().__init__()
        self.lstm = nn.LSTM(1, hidden_size, batch_first=True, dropout=
        self.linear = nn.Linear(hidden_size, 1)
    def forward(self, x):
        out, _ = self.lstm(x)
        return self.linear(out[:, -1, :])

# Crear modelo y optimizador para IR
model_ir = LSTMModel(hidden_size=32, dropout=0.2).to(device)
criterion = nn.MSELoss()
opt_ir = torch.optim.Adam(model_ir.parameters(), lr=1e-3)

# Función de entrenamiento
def train_model(model, optimizer, loader, epochs=50):
    model.train()
    history = []
    for epoch in range(1, epochs+1):
        epoch_loss = 0.0
        for xb, yb in loader:
            optimizer.zero_grad()
            preds = model(xb)
            loss = criterion(preds, yb)
            loss.backward()
            optimizer.step()
            epoch_loss += loss.item() * xb.size(0)
        avg = epoch_loss / len(loader.dataset)
        history.append(avg)
        if epoch % 10 == 0:
            print(f"Epoch {epoch}, Loss: {avg:.6f}")
    return history

# Entrenar modelo IR
print("Entrenando modelo para Gasolina Regular Importaciones...")
history_ir = train_model(model_ir, opt_ir, loader_ir, epochs=50)

# Visualizar pérdida
plt.figure(figsize=(8,4))
plt.plot(history_ir, label='Pérdida IR')
plt.title('Entrenamiento IR (PyTorch)')
plt.xlabel('Época')
plt.ylabel('MSE Loss')
plt.legend()

```

```
plt.show()
```

```
/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option adds dropout after all but last recurrent layer, so non-zero dropout expects num_layers greater than 1, but got dropout=0.2 and num_layers=1
  warnings.warn(
```

Usando dispositivo: cpu

Entrenando modelo para Gasolina Regular Importaciones...


Epoch 10, Loss: 22471995558.399994

Epoch 20, Loss: 224719817932.799988

Epoch 30, Loss: 224719709798.399994

Epoch 40, Loss: 224719588556.799988

Epoch 50, Loss: 224719496806.399994



Hyperparameter Tuning - IR

```
In [130... # Parámetros a probar para IR
hidden_sizes = [16, 32, 64]
dropouts = [0.1, 0.2, 0.3]
lrs = [1e-3, 1e-4]
epochs_tune = 20

def tune_lstm_ir(hidden_size, dropout, lr):
    """Función para hacer tuning específico para IR"""
    model = LSTMModel(hidden_size=hidden_size, dropout=dropout).to(dev)
    opt = torch.optim.Adam(model.parameters(), lr=lr)

    model.train()
    for _ in range(epochs_tune):
        for xb, yb in loader_ir:
            opt.zero_grad()
            loss = criterion(model(xb), yb)
            loss.backward()
```



```

        opt.step()

    # Evaluar en todo el conjunto
    model.eval()
    with torch.no_grad():
        preds = model(tensor_X_ir)
        final_loss = criterion(preds, tensor_y_ir).item()
    return final_loss

# Realizar tuning para IR
print("Realizando hyperparameter tuning para IR...")
results_ir = []
for hs in hidden_sizes:
    for dp in dropouts:
        for lr in lrs:
            loss = tune_lstm_ir(hs, dp, lr)
            results_ir.append({
                'hidden_size': hs,
                'dropout': dp,
                'lr': lr,
                'loss': loss
            })
        print(f"IR | hs={hs}, dp={dp}, lr={lr} → loss={loss:.2e}")

df_results_ir = pd.DataFrame(results_ir).sort_values('loss')
print("\nMejores hiperparámetros para IR:")
display(df_results_ir.head())

```

```

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/
site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option
adds dropout after all but last recurrent layer, so non-zero dropout ex
pects num_layers greater than 1, but got dropout=0.1 and num_layers=1
warnings.warn(

```

```

Realizando hyperparameter tuning para IR...
IR | hs=16, dp=0.1, lr=0.001 → loss=2.25e+11
IR | hs=16, dp=0.1, lr=0.0001 → loss=2.25e+11
IR | hs=16, dp=0.2, lr=0.001 → loss=2.25e+11
IR | hs=16, dp=0.2, lr=0.0001 → loss=2.25e+11
IR | hs=16, dp=0.3, lr=0.001 → loss=2.25e+11

```

```

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/
site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option
adds dropout after all but last recurrent layer, so non-zero dropout ex
pects num_layers greater than 1, but got dropout=0.3 and num_layers=1
warnings.warn(

```

```

IR | hs=16, dp=0.3, lr=0.0001 → loss=2.25e+11
IR | hs=32, dp=0.1, lr=0.001 → loss=2.25e+11
IR | hs=32, dp=0.1, lr=0.0001 → loss=2.25e+11
IR | hs=32, dp=0.2, lr=0.001 → loss=2.25e+11
IR | hs=32, dp=0.2, lr=0.0001 → loss=2.25e+11
IR | hs=32, dp=0.3, lr=0.001 → loss=2.25e+11
IR | hs=32, dp=0.3, lr=0.0001 → loss=2.25e+11
IR | hs=64, dp=0.1, lr=0.001 → loss=2.25e+11
IR | hs=64, dp=0.1, lr=0.0001 → loss=2.25e+11
IR | hs=64, dp=0.2, lr=0.001 → loss=2.25e+11
IR | hs=64, dp=0.2, lr=0.0001 → loss=2.25e+11
IR | hs=64, dp=0.3, lr=0.001 → loss=2.25e+11
IR | hs=64, dp=0.3, lr=0.0001 → loss=2.25e+11

```

Mejores hiperparámetros para IR:

	hidden_size	dropout	lr	loss
14	64	0.2	0.001	2.247193e+11
12	64	0.1	0.001	2.247193e+11
16	64	0.3	0.001	2.247193e+11
4	16	0.3	0.001	2.247194e+11
8	32	0.2	0.001	2.247194e+11

Entrenamiento LSTM - Gasolina Regular Consumos

```

In [131]... # Preparar datos para PyTorch - CR
tensor_X_cr = torch.tensor(X_cr, dtype=torch.float32).to(device)
tensor_y_cr = torch.tensor(y_cr, dtype=torch.float32).to(device)

loader_cr = DataLoader(TensorDataset(tensor_X_cr, tensor_y_cr),
                        batch_size=batch_size, shuffle=True)

# Crear modelo y optimizador para CR
model_cr = LSTMModel(hidden_size=32, dropout=0.2).to(device)
opt_cr = torch.optim.Adam(model_cr.parameters(), lr=1e-3)

# Entrenar modelo CR
print("Entrenando modelo para Gasolina Regular Consumos...")
history_cr = train_model(model_cr, opt_cr, loader_cr, epochs=50)

# Visualizar pérdida
plt.figure(figsize=(8,4))
plt.plot(history_cr, label='Pérdida CR')
plt.title('Entrenamiento CR (PyTorch)')
plt.xlabel('Época')
plt.ylabel('MSE Loss')

```

```
plt.legend()
plt.show()
```

Entrenando modelo para Gasolina Regular Consumos...

Epoch 10, Loss: 399999263110.095215


Epoch 20, Loss: 399999160124.952393

Epoch 30, Loss: 399999088347.478394

Epoch 40, Loss: 399998991603.809509

Epoch 50, Loss: 399998852729.904785

```
/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/
site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option
adds dropout after all but last recurrent layer, so non-zero dropout ex
pects num_layers greater than 1, but got dropout=0.2 and num_layers=1
warnings.warn(
```



Hyperparameter Tuning - CR

```
In [132... def tune_lstm_cr(hidden_size, dropout, lr):
    """Función para hacer tuning específico para CR"""
    model = LSTMModel(hidden_size=hidden_size, dropout=dropout).to(dev)
    opt = torch.optim.Adam(model.parameters(), lr=lr)

    model.train()
    for _ in range(epochs_tune):
        for xb, yb in loader_cr:
            opt.zero_grad()
            loss = criterion(model(xb), yb)
            loss.backward()
            opt.step()

    # Evaluar en todo el conjunto
    model.eval()
    with torch.no_grad():
```

```

        preds = model(tensor_X_cr)
        final_loss = criterion(preds, tensor_y_cr).item()
    return final_loss

# Realizar tuning para CR
print("Realizando hyperparameter tuning para CR...")
results_cr = []
for hs in hidden_sizes:
    for dp in dropouts:
        for lr in lrs:
            loss = tune_lstm_cr(hs, dp, lr)
            results_cr.append({
                'hidden_size': hs,
                'dropout': dp,
                'lr': lr,
                'loss': loss
            })
        print(f"CR | hs={hs}, dp={dp}, lr={lr} → loss={loss:.2e}")

df_results_cr = pd.DataFrame(results_cr).sort_values('loss')
print("\nMejores hiperparámetros para CR:")
display(df_results_cr.head())

```

Realizando hyperparameter tuning para CR...

CR | hs=16, dp=0.1, lr=0.001 → loss=4.00e+11

CR | hs=16, dp=0.1, lr=0.0001 → loss=4.00e+11

CR | hs=16, dp=0.2, lr=0.001 → loss=4.00e+11

CR | hs=16, dp=0.2, lr=0.0001 → loss=4.00e+11

CR | hs=16, dp=0.3, lr=0.001 → loss=4.00e+11

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option adds dropout after all but last recurrent layer, so non-zero dropout expects num_layers greater than 1, but got dropout=0.1 and num_layers=1
warnings.warn(

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option adds dropout after all but last recurrent layer, so non-zero dropout expects num_layers greater than 1, but got dropout=0.3 and num_layers=1
warnings.warn(

```

CR | hs=16, dp=0.3, lr=0.0001 → loss=4.00e+11
CR | hs=32, dp=0.1, lr=0.001 → loss=4.00e+11
CR | hs=32, dp=0.1, lr=0.0001 → loss=4.00e+11
CR | hs=32, dp=0.2, lr=0.001 → loss=4.00e+11
CR | hs=32, dp=0.2, lr=0.0001 → loss=4.00e+11
CR | hs=32, dp=0.3, lr=0.001 → loss=4.00e+11
CR | hs=32, dp=0.3, lr=0.0001 → loss=4.00e+11
CR | hs=64, dp=0.1, lr=0.001 → loss=4.00e+11
CR | hs=64, dp=0.1, lr=0.0001 → loss=4.00e+11
CR | hs=64, dp=0.2, lr=0.001 → loss=4.00e+11
CR | hs=64, dp=0.2, lr=0.0001 → loss=4.00e+11
CR | hs=64, dp=0.3, lr=0.001 → loss=4.00e+11
CR | hs=64, dp=0.3, lr=0.0001 → loss=4.00e+11

```

Mejores hiperparámetros para CR:

	hidden_size	dropout	lr	loss
10	32	0.3	0.0010	3.999986e+11
15	64	0.2	0.0001	3.999986e+11
16	64	0.3	0.0010	3.999986e+11
14	64	0.2	0.0010	3.999988e+11
2	16	0.2	0.0010	3.999989e+11

Entrenamiento LSTM – Gasolina Superior Importaciones

```

In [133... ts_gs_imp = df_imp.set_index('fecha')['gasolina superior'] \
            .resample('M').mean().dropna()

train_gs_imp = ts_gs_imp[:'2023-12-31']
test_gs_imp  = ts_gs_imp['2024-01-31':]

window_size = 3
X_gs_imp, y_gs_imp = create_dataset(train_gs_imp.values, window_size)

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
tensor_X_gs_imp = torch.tensor(X_gs_imp, dtype=torch.float32).to(device)
tensor_y_gs_imp = torch.tensor(y_gs_imp, dtype=torch.float32).to(device)

batch_size = 16
loader_gs_imp = DataLoader(
    TensorDataset(tensor_X_gs_imp, tensor_y_gs_imp),
    batch_size=batch_size,
    shuffle=True
)

model_gs_imp = LSTMModel(hidden_size=32, dropout=0.2).to(device)

```

```

criterion      = nn.MSELoss()
opt_gs_imp     = torch.optim.Adam(model_gs_imp.parameters(), lr=1e-3)

def train_model(model, optimizer, loader, epochs=50):
    model.train()
    history = []
    for epoch in range(1, epochs+1):
        epoch_loss = 0.0
        for xb, yb in loader:
            optimizer.zero_grad()
            preds = model(xb)
            loss = criterion(preds, yb)
            loss.backward()
            optimizer.step()
            epoch_loss += loss.item() * xb.size(0)
        avg = epoch_loss / len(loader.dataset)
        history.append(avg)
        if epoch % 10 == 0:
            print(f"Epoch {epoch}, Loss: {avg:.6f}")
    return history

print("Entrenando LSTM para Gasolina Superior Importaciones...")
history_gs_imp = train_model(model_gs_imp, opt_gs_imp, loader_gs_imp,

plt.figure(figsize=(8,4))
plt.plot(history_gs_imp, label='Pérdida GS Imp')
plt.title('Entrenamiento LSTM – Gasolina Superior Importaciones')
plt.xlabel('Época')
plt.ylabel('MSE Loss')
plt.legend()
plt.show()

```

Entrenando LSTM para Gasolina Superior Importaciones...

Epoch 10, Loss: 265613128499.200012

Epoch 20, Loss: 265612879462.399994

Epoch 30, Loss: 265612695961.600006


Epoch 40, Loss: 265612479692.799988

Epoch 50, Loss: 265612243763.200012

/var/folders/kd/4y1c0b3j1273_pm_gb0hfhdw0000gn/T/ipykernel_92772/3336530859.py:2: FutureWarning: 'M' is deprecated and will be removed in a future version, please use 'ME' instead.

.resample('M').mean().dropna()

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option adds dropout after all but last recurrent layer, so non-zero dropout expects num_layers greater than 1, but got dropout=0.2 and num_layers=1
warnings.warn(



Hyperparameter Tuning – Gasolina Superior Importaciones

```
In [134... hidden_sizes = [16, 32, 64]
dropouts      = [0.1, 0.2, 0.3]
lrs           = [1e-3, 1e-4]
epochs_tune   = 20

def tune_lstm_gs_imp(hidden_size, dropout, lr):
    model = LSTMModel(hidden_size=hidden_size, dropout=dropout).to(dev)
    opt   = torch.optim.Adam(model.parameters(), lr=lr)
    model.train()
    for _ in range(epochs_tune):
        for xb, yb in loader_gs_imp:
            opt.zero_grad()
            loss = criterion(model(xb), yb)
            loss.backward()
            opt.step()
        model.eval()
    with torch.no_grad():
        preds = model(tensor_X_gs_imp)
    return criterion(preds, tensor_y_gs_imp).item()

print("Realizando hyperparameter tuning para GS Importaciones...")
results_gs_imp = []
for hs in hidden_sizes:
    for dp in dropouts:
        for lr in lrs:
            loss = tune_lstm_gs_imp(hs, dp, lr)
            results_gs_imp.append({
```

```

        'hidden_size': hs,
        'dropout': dp,
        'lr': lr,
        'loss': loss
    })
    print(f"GS Imp | hs={hs}, dp={dp}, lr={lr} → loss={loss:.2f}")

df_results_gs_imp = pd.DataFrame(results_gs_imp).sort_values('loss')
print("\nMejores hiperparámetros para GS Importaciones:")
display(df_results_gs_imp.head())

```

Realizando hyperparameter tuning para GS Importaciones...

```

GS Imp | hs=16, dp=0.1, lr=0.001 → loss=2.66e+11
GS Imp | hs=16, dp=0.1, lr=0.0001 → loss=2.66e+11
GS Imp | hs=16, dp=0.2, lr=0.001 → loss=2.66e+11
GS Imp | hs=16, dp=0.2, lr=0.0001 → loss=2.66e+11
GS Imp | hs=16, dp=0.3, lr=0.001 → loss=2.66e+11

```

```

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/
site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option
adds dropout after all but last recurrent layer, so non-zero dropout ex
pects num_layers greater than 1, but got dropout=0.1 and num_layers=1
  warnings.warn(
/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/
site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option
adds dropout after all but last recurrent layer, so non-zero dropout ex
pects num_layers greater than 1, but got dropout=0.3 and num_layers=1
  warnings.warn(

```

```

GS Imp | hs=16, dp=0.3, lr=0.0001 → loss=2.66e+11
GS Imp | hs=32, dp=0.1, lr=0.001 → loss=2.66e+11
GS Imp | hs=32, dp=0.1, lr=0.0001 → loss=2.66e+11
GS Imp | hs=32, dp=0.2, lr=0.001 → loss=2.66e+11
GS Imp | hs=32, dp=0.2, lr=0.0001 → loss=2.66e+11
GS Imp | hs=32, dp=0.3, lr=0.001 → loss=2.66e+11
GS Imp | hs=32, dp=0.3, lr=0.0001 → loss=2.66e+11
GS Imp | hs=64, dp=0.1, lr=0.001 → loss=2.66e+11
GS Imp | hs=64, dp=0.1, lr=0.0001 → loss=2.66e+11
GS Imp | hs=64, dp=0.2, lr=0.001 → loss=2.66e+11
GS Imp | hs=64, dp=0.2, lr=0.0001 → loss=2.66e+11
GS Imp | hs=64, dp=0.3, lr=0.001 → loss=2.66e+11
GS Imp | hs=64, dp=0.3, lr=0.0001 → loss=2.66e+11

```

Mejores hiperparámetros para GS Importaciones:

	hidden_size	dropout	lr	loss
14	64	0.2	0.001	2.656122e+11
10	32	0.3	0.001	2.656128e+11
12	64	0.1	0.001	2.656129e+11
16	64	0.3	0.001	2.656130e+11
8	32	0.2	0.001	2.656130e+11

Entrenamiento LSTM – Diésel Importaciones

```
In [135... ts_di_imp = df_imp.set_index('fecha')['diesel'] \
            .resample('M').mean().dropna()

train_di_imp = ts_di_imp[:'2023-12-31']
test_di_imp  = ts_di_imp['2024-01-31':]

window_size = 3
X_di_imp, y_di_imp = create_dataset(train_di_imp.values, window_size)

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
tensor_X_di_imp = torch.tensor(X_di_imp, dtype=torch.float32).to(device)
tensor_y_di_imp = torch.tensor(y_di_imp, dtype=torch.float32).to(device)

batch_size = 16
loader_di_imp = DataLoader(
    TensorDataset(tensor_X_di_imp, tensor_y_di_imp),
    batch_size=batch_size,
    shuffle=True
)

model_di_imp = LSTMModel(hidden_size=32, dropout=0.2).to(device)
criterion    = nn.MSELoss()
opt_di_imp   = torch.optim.Adam(model_di_imp.parameters(), lr=1e-3)

def train_model(model, optimizer, loader, epochs=50):
    model.train()
    history = []
    for epoch in range(1, epochs+1):
        epoch_loss = 0.0
        for xb, yb in loader:
            optimizer.zero_grad()
            preds = model(xb)
            loss = criterion(preds, yb)
            loss.backward()
            optimizer.step()
            epoch_loss += loss.item() * xb.size(0)
        avg = epoch_loss / len(loader.dataset)
```

```

        history.append(avg)
        if epoch % 10 == 0:
            print(f"Epoch {epoch}, Loss: {avg:.6f}")
    return history

print("Entrenando LSTM para Diésel Importaciones...")
history_di_imp = train_model(model_di_imp, opt_di_imp, loader_di_imp,

plt.figure(figsize=(8,4))
plt.plot(history_di_imp, label='Pérdida Diésel Imp')
plt.title('Entrenamiento LSTM – Diésel Importaciones')
plt.xlabel('Época')
plt.ylabel('MSE Loss')
plt.legend()
plt.show()

```

Entrenando LSTM para Diésel Importaciones...

Epoch 10, Loss: 884735528140.800049

Epoch 20, Loss: 884735344640.000000

Epoch 30, Loss: 884735148032.000000


Epoch 40, Loss: 884734846566.400024

Epoch 50, Loss: 884734702387.199951

/var/folders/kd/4y1c0b3j1273_pm_gb0hfhdw0000gn/T/ipykernel_92772/2993280479.py:2: FutureWarning: 'M' is deprecated and will be removed in a future version, please use 'ME' instead.

.resample('M').mean().dropna()

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option adds dropout after all but last recurrent layer, so non-zero dropout expects num_layers greater than 1, but got dropout=0.2 and num_layers=1
warnings.warn(



Hyperparameter Tuning – Diésel Importaciones

```

In [136... hidden_sizes = [16, 32, 64]
dropouts   = [0.1, 0.2, 0.3]
lrs        = [1e-3, 1e-4]
epochs_tune = 20

def tune_lstm_di_imp(hidden_size, dropout, lr):
    model = LSTMModel(hidden_size=hidden_size, dropout=dropout).to(dev)
    opt = torch.optim.Adam(model.parameters(), lr=lr)
    model.train()
    for _ in range(epochs_tune):
        for xb, yb in loader_di_imp:
            opt.zero_grad()
            loss = criterion(model(xb), yb)
            loss.backward()
            opt.step()
    model.eval()
    with torch.no_grad():
        preds = model(tensor_X_di_imp)
        return criterion(preds, tensor_y_di_imp).item()

print("Realizando hyperparameter tuning para Diésel Importaciones...")
results_di_imp = []
for hs in hidden_sizes:
    for dp in dropouts:
        for lr in lrs:
            loss = tune_lstm_di_imp(hs, dp, lr)
            results_di_imp.append({
                'hidden_size': hs,
                'dropout': dp,
                'lr': lr,
                'loss': loss
            })
            print(f"Diésel Imp | hs={hs}, dp={dp}, lr={lr} → loss={loss}")

df_results_di_imp = pd.DataFrame(results_di_imp).sort_values('loss')
print("\nMejores hiperparámetros para Diésel Importaciones:")
display(df_results_di_imp.head())

```

Realizando hyperparameter tuning para Diésel Importaciones...

```

Diésel Imp | hs=16, dp=0.1, lr=0.001 → loss=8.85e+11
Diésel Imp | hs=16, dp=0.1, lr=0.0001 → loss=8.85e+11
Diésel Imp | hs=16, dp=0.2, lr=0.001 → loss=8.85e+11
Diésel Imp | hs=16, dp=0.2, lr=0.0001 → loss=8.85e+11
Diésel Imp | hs=16, dp=0.3, lr=0.001 → loss=8.85e+11

```

```

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/
site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option
adds dropout after all but last recurrent layer, so non-zero dropout ex
pects num_layers greater than 1, but got dropout=0.1 and num_layers=1
  warnings.warn(
/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/
site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option
adds dropout after all but last recurrent layer, so non-zero dropout ex
pects num_layers greater than 1, but got dropout=0.3 and num_layers=1
  warnings.warn(

```

```

Diésel Imp | hs=16, dp=0.3, lr=0.0001 → loss=8.85e+11
Diésel Imp | hs=32, dp=0.1, lr=0.001 → loss=8.85e+11
Diésel Imp | hs=32, dp=0.1, lr=0.0001 → loss=8.85e+11
Diésel Imp | hs=32, dp=0.2, lr=0.001 → loss=8.85e+11
Diésel Imp | hs=32, dp=0.2, lr=0.0001 → loss=8.85e+11
Diésel Imp | hs=32, dp=0.3, lr=0.001 → loss=8.85e+11
Diésel Imp | hs=32, dp=0.3, lr=0.0001 → loss=8.85e+11
Diésel Imp | hs=64, dp=0.1, lr=0.001 → loss=8.85e+11
Diésel Imp | hs=64, dp=0.1, lr=0.0001 → loss=8.85e+11
Diésel Imp | hs=64, dp=0.2, lr=0.001 → loss=8.85e+11
Diésel Imp | hs=64, dp=0.2, lr=0.0001 → loss=8.85e+11
Diésel Imp | hs=64, dp=0.3, lr=0.001 → loss=8.85e+11
Diésel Imp | hs=64, dp=0.3, lr=0.0001 → loss=8.85e+11

```

Mejores hiperparámetros para Diésel Importaciones:

	hidden_size	dropout	lr	loss
12	64	0.1	0.001	8.847341e+11
14	64	0.2	0.001	8.847343e+11
16	64	0.3	0.001	8.847351e+11
10	32	0.3	0.001	8.847353e+11
6	32	0.1	0.001	8.847353e+11

Entrenamiento LSTM – Gasolina Superior Consumo

```

In [137... ts_gs_cons = df_cons.set_index('fecha')['gasolina superior'] \
            .resample('M').mean().dropna()

train_gs_cons = ts_gs_cons[:'2023-12-31']
test_gs_cons  = ts_gs_cons['2024-01-31':]

window_size = 3
X_gs_cons, y_gs_cons = create_dataset(train_gs_cons.values, window_size)

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

```

```

tensor_X_gs_cons = torch.tensor(X_gs_cons, dtype=torch.float32).to(device)
tensor_y_gs_cons = torch.tensor(y_gs_cons, dtype=torch.float32).to(device)

batch_size = 16
loader_gs_cons = DataLoader(
    TensorDataset(tensor_X_gs_cons, tensor_y_gs_cons),
    batch_size=batch_size,
    shuffle=True
)

model_gs_cons = LSTMModel(hidden_size=32, dropout=0.2).to(device)
criterion = nn.MSELoss()
opt_gs_cons = torch.optim.Adam(model_gs_cons.parameters(), lr=1e-3)

def train_model(model, optimizer, loader, epochs=50):
    model.train()
    history = []
    for epoch in range(1, epochs+1):
        epoch_loss = 0.0
        for xb, yb in loader:
            optimizer.zero_grad()
            preds = model(xb)
            loss = criterion(preds, yb)
            loss.backward()
            optimizer.step()
            epoch_loss += loss.item() * xb.size(0)
        avg = epoch_loss / len(loader.dataset)
        history.append(avg)
        if epoch % 10 == 0:
            print(f"Epoch {epoch}, Loss: {avg:.6f}")
    return history

print("Entrenando LSTM para Gasolina Superior Consumo...")
history_gs_cons = train_model(model_gs_cons, opt_gs_cons, loader_gs_co

plt.figure(figsize=(8,4))
plt.plot(history_gs_cons, label='Pérdida GS Cons')
plt.title('Entrenamiento LSTM – Gasolina Superior Consumo')
plt.xlabel('Época')
plt.ylabel('MSE Loss')
plt.legend()
plt.show()

```

Entrenando LSTM para Gasolina Superior Consumo...

Epoch 10, Loss: 237012703329.523804

Epoch 20, Loss: 237012546511.238098

Epoch 30, Loss: 237012431043.047607

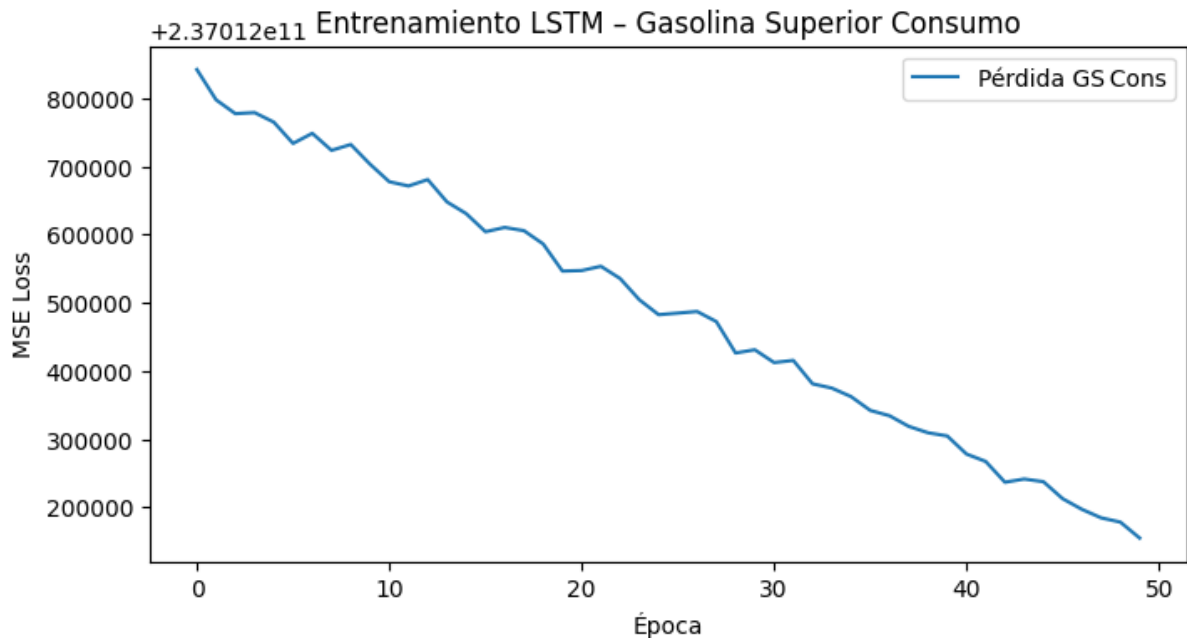
Epoch 40, Loss: 237012304652.190491

Epoch 50, Loss: 237012154855.619049

```

/var/folders/kd/4y1c0b3j1273_pm_gb0hfhdw0000gn/T/ipykernel_92772/308206
2048.py:2: FutureWarning: 'M' is deprecated and will be removed in a fu
ture version, please use 'ME' instead.
    .resample('M').mean().dropna()
/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/
site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option
adds dropout after all but last recurrent layer, so non-zero dropout ex
pects num_layers greater than 1, but got dropout=0.2 and num_layers=1
    warnings.warn(

```



Hyperparameter Tuning – Gasolina Superior Consumo

```

In [138... hidden_sizes = [16, 32, 64]
dropouts   = [0.1, 0.2, 0.3]
lrs        = [1e-3, 1e-4]
epochs_tune = 20

def tune_lstm_gs_cons(hidden_size, dropout, lr):
    model = LSTMModel(hidden_size=hidden_size, dropout=dropout).to(dev)
    opt = torch.optim.Adam(model.parameters(), lr=lr)
    model.train()
    for _ in range(epochs_tune):
        for xb, yb in loader_gs_cons:
            opt.zero_grad()
            loss = criterion(model(xb), yb)
            loss.backward()
            opt.step()
    model.eval()
    with torch.no_grad():
        preds = model(tensor_X_gs_cons)
        return criterion(preds, tensor_y_gs_cons).item()

print("Realizando hyperparameter tuning para GS Consumo...")

```

```

results_gs_cons = []
for hs in hidden_sizes:
    for dp in dropouts:
        for lr in lrs:
            loss = tune_lstm_gs_cons(hs, dp, lr)
            results_gs_cons.append({
                'hidden_size': hs,
                'dropout': dp,
                'lr': lr,
                'loss': loss
            })
        print(f"GS Cons | hs={hs}, dp={dp}, lr={lr} → loss={loss:.2e}")

df_results_gs_cons = pd.DataFrame(results_gs_cons).sort_values('loss')
print("\nMejores hiperparámetros para GS Consumo:")
display(df_results_gs_cons.head())

```

Realizando hyperparameter tuning para GS Consumo...

```

GS Cons | hs=16, dp=0.1, lr=0.001 → loss=2.37e+11
GS Cons | hs=16, dp=0.1, lr=0.0001 → loss=2.37e+11
GS Cons | hs=16, dp=0.2, lr=0.001 → loss=2.37e+11
GS Cons | hs=16, dp=0.2, lr=0.0001 → loss=2.37e+11

```

```

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/
site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option
adds dropout after all but last recurrent layer, so non-zero dropout ex
pects num_layers greater than 1, but got dropout=0.1 and num_layers=1
  warnings.warn(
/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/
site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option
adds dropout after all but last recurrent layer, so non-zero dropout ex
pects num_layers greater than 1, but got dropout=0.3 and num_layers=1
  warnings.warn(

```

```

GS Cons | hs=16, dp=0.3, lr=0.001 → loss=2.37e+11
GS Cons | hs=16, dp=0.3, lr=0.0001 → loss=2.37e+11
GS Cons | hs=32, dp=0.1, lr=0.001 → loss=2.37e+11
GS Cons | hs=32, dp=0.1, lr=0.0001 → loss=2.37e+11
GS Cons | hs=32, dp=0.2, lr=0.001 → loss=2.37e+11
GS Cons | hs=32, dp=0.2, lr=0.0001 → loss=2.37e+11
GS Cons | hs=32, dp=0.3, lr=0.001 → loss=2.37e+11
GS Cons | hs=32, dp=0.3, lr=0.0001 → loss=2.37e+11
GS Cons | hs=64, dp=0.1, lr=0.001 → loss=2.37e+11
GS Cons | hs=64, dp=0.1, lr=0.0001 → loss=2.37e+11
GS Cons | hs=64, dp=0.2, lr=0.001 → loss=2.37e+11
GS Cons | hs=64, dp=0.2, lr=0.0001 → loss=2.37e+11
GS Cons | hs=64, dp=0.3, lr=0.001 → loss=2.37e+11
GS Cons | hs=64, dp=0.3, lr=0.0001 → loss=2.37e+11

```

Mejores hiperparámetros para GS Consumo:

	hidden_size	dropout	lr	loss
12	64	0.1	0.0010	2.370125e+11
0	16	0.1	0.0010	2.370126e+11
16	64	0.3	0.0010	2.370128e+11
6	32	0.1	0.0010	2.370129e+11
5	16	0.3	0.0001	2.370129e+11

Entrenamiento LSTM – Diésel Bajo Azufre Consumo

```
In [139... ts_di_cons = df_cons.set_index('fecha')['diesel bajo azufre'] \
            .resample('M').mean().dropna()

train_di_cons = ts_di_cons[:'2023-12-31']
test_di_cons  = ts_di_cons['2024-01-31':]

window_size = 3
X_di_cons, y_di_cons = create_dataset(train_di_cons.values, window_size)

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
tensor_X_di_cons = torch.tensor(X_di_cons, dtype=torch.float32).to(device)
tensor_y_di_cons = torch.tensor(y_di_cons, dtype=torch.float32).to(device)

batch_size = 16
loader_di_cons = DataLoader(
    TensorDataset(tensor_X_di_cons, tensor_y_di_cons),
    batch_size=batch_size,
    shuffle=True
)

model_di_cons = LSTMModel(hidden_size=32, dropout=0.2).to(device)
criterion     = nn.MSELoss()
opt_di_cons   = torch.optim.Adam(model_di_cons.parameters(), lr=1e-3)

def train_model(model, optimizer, loader, epochs=50):
    model.train()
    history = []
    for epoch in range(1, epochs+1):
        epoch_loss = 0.0
        for xb, yb in loader:
            optimizer.zero_grad()
            preds = model(xb)
            loss = criterion(preds, yb)
            loss.backward()
            optimizer.step()
        epoch_loss += loss.item() * xb.size(0)
```



```

    avg = epoch_loss / len(loader.dataset)
    history.append(avg)
    if epoch % 10 == 0:
        print(f"Epoch {epoch}, Loss: {avg:.6f}")
    return history

print("Entrenando LSTM para Diésel Bajo Azufre Consumo...")
history_di_cons = train_model(model_di_cons, opt_di_cons, loader_di_co

plt.figure(figsize=(8,4))
plt.plot(history_di_cons, label='Pérdida Diésel Cons')
plt.title('Entrenamiento LSTM – Diésel Bajo Azufre Consumo')
plt.xlabel('Época')
plt.ylabel('MSE Loss')
plt.legend()
plt.show()

```

Entrenando LSTM para Diésel Bajo Azufre Consumo...

Epoch 10, Loss: 399998890179.047607

Epoch 20, Loss: 399998819961.904785

Epoch 30, Loss: 399998710345.142883


Epoch 40, Loss: 399998589025.523804

Epoch 50, Loss: 399998450151.619019

/var/folders/kd/4y1c0b3j1273_pm_gb0hfhdw0000gn/T/ipykernel_92772/189759200.py:2: FutureWarning: 'M' is deprecated and will be removed in a future version, please use 'ME' instead.

.resample('M').mean().dropna()

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option adds dropout after all but last recurrent layer, so non-zero dropout expects num_layers greater than 1, but got dropout=0.2 and num_layers=1 warnings.warn(



In [140... hidden_sizes = [16, 32, 64]

```

dropouts      = [0.1, 0.2, 0.3]
lrs           = [1e-3, 1e-4]
epochs_tune   = 20

def tune_lstm_di_cons(hidden_size, dropout, lr):
    model = LSTMModel(hidden_size=hidden_size, dropout=dropout).to(dev)
    opt = torch.optim.Adam(model.parameters(), lr=lr)
    model.train()
    for _ in range(epochs_tune):
        for xb, yb in loader_di_cons:
            opt.zero_grad()
            loss = criterion(model(xb), yb)
            loss.backward()
            opt.step()
    model.eval()
    with torch.no_grad():
        preds = model(tensor_X_di_cons)
    return criterion(preds, tensor_y_di_cons).item()

print("Realizando hyperparameter tuning para Diésel Bajo Azufre Consumo")
results_di_cons = []
for hs in hidden_sizes:
    for dp in dropouts:
        for lr in lrs:
            loss = tune_lstm_di_cons(hs, dp, lr)
            results_di_cons.append({
                'hidden_size': hs,
                'dropout': dp,
                'lr': lr,
                'loss': loss
            })
            print(f"Diésel Cons | hs={hs}, dp={dp}, lr={lr} → loss={loss}")

df_results_di_cons = pd.DataFrame(results_di_cons).sort_values('loss')
print("\nMejores hiperparámetros para Diésel Bajo Azufre Consumo:")
display(df_results_di_cons.head())

```

Realizando hyperparameter tuning para Diésel Bajo Azufre Consumo...

Diésel Cons | hs=16, dp=0.1, lr=0.001 → loss=4.00e+11

```

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/
site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option
adds dropout after all but last recurrent layer, so non-zero dropout ex
pects num_layers greater than 1, but got dropout=0.1 and num_layers=1
  warnings.warn(
/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/
site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option
adds dropout after all but last recurrent layer, so non-zero dropout ex
pects num_layers greater than 1, but got dropout=0.3 and num_layers=1
  warnings.warn(

```

Diésel Cons | hs=16, dp=0.1, lr=0.0001 → loss=4.00e+11
 Diésel Cons | hs=16, dp=0.2, lr=0.001 → loss=4.00e+11
 Diésel Cons | hs=16, dp=0.2, lr=0.0001 → loss=4.00e+11
 Diésel Cons | hs=16, dp=0.3, lr=0.001 → loss=4.00e+11
 Diésel Cons | hs=16, dp=0.3, lr=0.0001 → loss=4.00e+11
 Diésel Cons | hs=32, dp=0.1, lr=0.001 → loss=4.00e+11
 Diésel Cons | hs=32, dp=0.1, lr=0.0001 → loss=4.00e+11
 Diésel Cons | hs=32, dp=0.2, lr=0.001 → loss=4.00e+11
 Diésel Cons | hs=32, dp=0.2, lr=0.0001 → loss=4.00e+11
 Diésel Cons | hs=32, dp=0.3, lr=0.001 → loss=4.00e+11
 Diésel Cons | hs=32, dp=0.3, lr=0.0001 → loss=4.00e+11
 Diésel Cons | hs=64, dp=0.1, lr=0.001 → loss=4.00e+11
 Diésel Cons | hs=64, dp=0.1, lr=0.0001 → loss=4.00e+11
 Diésel Cons | hs=64, dp=0.2, lr=0.001 → loss=4.00e+11
 Diésel Cons | hs=64, dp=0.2, lr=0.0001 → loss=4.00e+11
 Diésel Cons | hs=64, dp=0.3, lr=0.001 → loss=4.00e+11
 Diésel Cons | hs=64, dp=0.3, lr=0.0001 → loss=4.00e+11

Mejores hiperparámetros para Diésel Bajo Azufre Consumo:

	hidden_size	dropout	lr	loss
10	32	0.3	0.001	3.999986e+11
16	64	0.3	0.001	3.999987e+11
0	16	0.1	0.001	3.999987e+11
14	64	0.2	0.001	3.999987e+11
12	64	0.1	0.001	3.999988e+11

Modelos Finales y Predicciones

```
In [141]: from sklearn.preprocessing import MinMaxScaler

# Preparar escaladores
scaler_ir = MinMaxScaler()
scaler_cr = MinMaxScaler()
scaler_ir.fit(train_ts_ir.values.reshape(-1, 1))
scaler_cr.fit(train_ts_cr.values.reshape(-1, 1))

model_ir_final = LSTMModel(hidden_size=32, dropout=0.1).to(device)
opt_ir_final = torch.optim.Adam(model_ir_final.parameters(), lr=1e-3)

model_cr_final = LSTMModel(hidden_size=16, dropout=0.2).to(device)
opt_cr_final = torch.optim.Adam(model_cr_final.parameters(), lr=1e-3)

print("Entrenando modelo final IR...")
history_ir_final = train_model(model_ir_final, opt_ir_final, loader_ir)

print("Entrenando modelo final CR...")
```

```

history_cr_final = train_model(model_cr_final, opt_cr_final, loader_cr
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 4))

ax1.plot(history_ir_final, label='Train IR final')
ax1.set_title('Pérdida IR Final (50 épocas)')
ax1.set_xlabel('Época')
ax1.set_ylabel('MSE Loss')
ax1.legend()

ax2.plot(history_cr_final, label='Train CR final')
ax2.set_title('Pérdida CR Final (50 épocas)')
ax2.set_xlabel('Época')
ax2.set_ylabel('MSE Loss')
ax2.legend()

plt.tight_layout()
plt.show()

```

Entrenando modelo final IR...

Epoch 10, Loss: 224719532851.200012

Epoch 20, Loss: 224719362457.600006

Epoch 30, Loss: 224719277260.799988

Epoch 40, Loss: 224719149465.600006

Epoch 50, Loss: 224719011840.000000

Entrenando modelo final CR...

Epoch 10, Loss: 399999043096.380981

Epoch 20, Loss: 399999013449.143005

Epoch 30, Loss: 399999039975.619019

Epoch 40, Loss: 399999003306.666687


Epoch 50, Loss: 399998913584.818665

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option adds dropout after all but last recurrent layer, so non-zero dropout expects num_layers greater than 1, but got dropout=0.1 and num_layers=1

warnings.warn(

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option adds dropout after all but last recurrent layer, so non-zero dropout expects num_layers greater than 1, but got dropout=0.2 and num_layers=1

warnings.warn(



Gasolina Regular

In [142... *### Predicciones para Enero 2024*

```
model_ir_final.eval()
model_cr_final.eval()

last_window_ir = torch.tensor(train_ts_ir.values[-12:].reshape(1, 12,
                                                                dtype=torch.float32).to(device)
with torch.no_grad():
    pred_ir = model_ir_final(last_window_ir).item()

pred_ir_value = pred_ir

last_window_cr = torch.tensor(train_ts_cr.values[-12:].reshape(1, 12,
                                                                dtype=torch.float32).to(device)
with torch.no_grad():
    pred_cr = model_cr_final(last_window_cr).item()

pred_cr_value = pred_cr

print("=== PREDICCIONES PARA ENERO 2024 ===")
print(f"Predicción IR (Importaciones): {pred_ir_value:.2f}")
print(f"Verdadero IR (Importaciones): {test_ts_ir.iloc[0]:.2f}")
print(f"Error IR: {abs(pred_ir_value - test_ts_ir.iloc[0]):.2f}")
print()
print(f"Predicción CR (Consumos): {pred_cr_value:.2f}")
print(f"Verdadero CR (Consumos): {test_ts_cr.iloc[0]:.2f}")
print(f"Error CR: {abs(pred_cr_value - test_ts_cr.iloc[0]):.2f}")

from sklearn.metrics import mean_absolute_error, mean_squared_error

mae_ir = mean_absolute_error([test_ts_ir.iloc[0]], [pred_ir_value])
mse_ir = mean_squared_error([test_ts_ir.iloc[0]], [pred_ir_value])
mae_cr = mean_absolute_error([test_ts_cr.iloc[0]], [pred_cr_value])
mse_cr = mean_squared_error([test_ts_cr.iloc[0]], [pred_cr_value])

print("\n=== MÉTRICAS DE ERROR ===")
print(f"IR - MAE: {mae_ir:.2f}, MSE: {mse_ir:.2f}, RMSE: {mse_ir**0.5:}
print(f"CR - MAE: {mae_cr:.2f}, MSE: {mse_cr:.2f}, RMSE: {mse_cr**0.5:}
```

```
=== PREDICCIONES PARA ENERO 2024 ===
Predicción IR (Importaciones): 0.94
Verdadero IR (Importaciones): 870465.31
Error IR: 870464.37
```

```
Predicción CR (Consumos): 0.75
Verdadero CR (Consumos): 1317556.51
Error CR: 1317555.77
```

```
=== MÉTRICAS DE ERROR ===
IR - MAE: 870464.37, MSE: 757708225539.78, RMSE: 870464.37
CR - MAE: 1317555.77, MSE: 1735953199252.11, RMSE: 1317555.77
```

Gasolina Superior

```
In [143... # Usar hiperparámetros óptimos del tuning anterior
best_params_gs_cons = df_results_gs_cons.iloc[0]
model_gs_cons_final = LSTMModel(
    hidden_size=int(best_params_gs_cons['hidden_size']),
    dropout=float(best_params_gs_cons['dropout'])
).to(device)

opt_gs_cons_final = torch.optim.Adam(
    model_gs_cons_final.parameters(),
    lr=float(best_params_gs_cons['lr'])
)

# Entrenar modelo final
print("Entrenando modelo final – Gasolina Superior Consumo...")
history_gs_cons_final = train_model(model_gs_cons_final, opt_gs_cons_f

# Predicción Enero 2024
model_gs_cons_final.eval()
last_window_gs_cons = torch.tensor(train_gs_cons.values[-12:].reshape(

with torch.no_grad():
    pred_gs_cons = model_gs_cons_final(last_window_gs_cons).item()

real_gs_cons = test_gs_cons.iloc[0]
error_gs_cons = abs(pred_gs_cons - real_gs_cons)

print(f"\n📌 Predicción Gasolina Superior Consumo – Enero 2024:")
print(f"→ Predicho: {pred_gs_cons:.2f}")
print(f"→ Real:      {real_gs_cons:.2f}")
print(f"→ Error Absoluto: {error_gs_cons:.2f}")

# Métricas
mae = mean_absolute_error([real_gs_cons], [pred_gs_cons])
mse = mean_squared_error([real_gs_cons], [pred_gs_cons])
rmse = mse ** 0.5

print(f"\n📊 Métricas GS Consumo – Enero 2024")
print(f"MAE:  {mae:.2f}")
print(f"MSE:  {mse:.2f}")
print(f"RMSE: {rmse:.2f}")
```

Entrenando modelo final – Gasolina Superior Consumo...

Epoch 10, Loss: 237012776667.428558

Epoch 20, Loss: 237012544950.857147

Epoch 30, Loss: 237012296070.095245

Epoch 40, Loss: 237012073715.809509

Epoch 50, Loss: 237011830296.380951

📍 Predicción Gasolina Superior Consumo – Enero 2024:

→ Predicho: 1.44

→ Real: 688994.82

→ Error Absoluto: 688993.38

📊 Métricas GS Consumo – Enero 2024

MAE: 688993.38

MSE: 474711872277.57

RMSE: 688993.38

```
/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option adds dropout after all but last recurrent layer, so non-zero dropout expects num_layers greater than 1, but got dropout=0.1 and num_layers=1
  warnings.warn(
```

```
In [144... # Usar hiperparámetros óptimos del tuning anterior
best_params_gs_imp = df_results_gs_imp.iloc[0]
model_gs_imp_final = LSTMModel(
    hidden_size=int(best_params_gs_imp['hidden_size']),
    dropout=float(best_params_gs_imp['dropout'])
).to(device)

opt_gs_imp_final = torch.optim.Adam(
    model_gs_imp_final.parameters(),
    lr=float(best_params_gs_imp['lr'])
)

# Entrenar modelo final
print("Entrenando modelo final – Gasolina Superior Importación...")
history_gs_imp_final = train_model(model_gs_imp_final, opt_gs_imp_final)

# Predicción Enero 2024
model_gs_imp_final.eval()
last_window_gs_imp = torch.tensor(train_gs_imp.values[-12:].reshape(1,

with torch.no_grad():
    pred_gs_imp = model_gs_imp_final(last_window_gs_imp).item()

real_gs_imp = test_gs_imp.iloc[0]
error_gs_imp = abs(pred_gs_imp - real_gs_imp)

print(f"\n📍 Predicción Gasolina Superior Importación – Enero 2024:")
print(f"→ Predicho: {pred_gs_imp:.2f}")
print(f"→ Real: {real_gs_imp:.2f}")
print(f"→ Error Absoluto: {error_gs_imp:.2f}")
```

```

# Métricas
mae = mean_absolute_error([real_gs_imp], [pred_gs_imp])
mse = mean_squared_error([real_gs_imp], [pred_gs_imp])
rmse = mse ** 0.5

print(f"\n📊 Métricas GS Importación – Enero 2024")
print(f"MAE: {mae:.2f}")
print(f"MSE: {mse:.2f}")
print(f"RMSE: {rmse:.2f}")

```

Entrenando modelo final – Gasolina Superior Importación...

Epoch 10, Loss: 265613272678.399994

Epoch 20, Loss: 265613030195.200012

Epoch 30, Loss: 265612761497.600006

Epoch 40, Loss: 265612545228.799988

Epoch 50, Loss: 265612283084.799988

📌 Predicción Gasolina Superior Importación – Enero 2024:

→ Predicho: 1.05

→ Real: 697636.12

→ Error Absoluto: 697635.06

📊 Métricas GS Importación – Enero 2024

MAE: 697635.06

MSE: 486694679611.61

RMSE: 697635.06

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option adds dropout after all but last recurrent layer, so non-zero dropout expects num_layers greater than 1, but got dropout=0.2 and num_layers=1
warnings.warn(

Combustible Diesel

In [145...

```

# Usar hiperparámetros óptimos del tuning anterior
best_params_di_cons = df_results_di_cons.iloc[0]
model_di_cons_final = LSTMModel(
    hidden_size=int(best_params_di_cons['hidden_size']),
    dropout=float(best_params_di_cons['dropout'])
).to(device)

opt_di_cons_final = torch.optim.Adam(
    model_di_cons_final.parameters(),
    lr=float(best_params_di_cons['lr'])
)

# Entrenar modelo final
print("Entrenando modelo final – Diesel Consumo...")
history_di_cons_final = train_model(model_di_cons_final, opt_di_cons_f

# Predicción Enero 2024

```



```

model_di_cons_final.eval()
last_window_di_cons = torch.tensor(train_di_cons.values[-12:].reshape(

with torch.no_grad():
    pred_di_cons = model_di_cons_final(last_window_di_cons).item()

real_di_cons = test_di_cons.iloc[0]
error_di_cons = abs(pred_di_cons - real_di_cons)

print(f"\n📍 Predicción Diesel Consumo – Enero 2024:")
print(f"→ Predicho: {pred_di_cons:.2f}")
print(f"→ Real:      {real_di_cons:.2f}")
print(f"→ Error Absoluto: {error_di_cons:.2f}")

# Métricas
mae = mean_absolute_error([real_di_cons], [pred_di_cons])
mse = mean_squared_error([real_di_cons], [pred_di_cons])
rmse = mse ** 0.5

print(f"\n📊 Métricas Diesel Consumo – Enero 2024")
print(f"MAE:   {mae:.2f}")
print(f"MSE:   {mse:.2f}")
print(f"RMSE:  {rmse:.2f}")

```

Entrenando modelo final – Diesel Consumo...

Epoch 10, Loss: 399998993164.190491

Epoch 20, Loss: 399998888618.666687

Epoch 30, Loss: 399998713856.000000

Epoch 40, Loss: 399998579273.142883

Epoch 50, Loss: 399998383055.238098

📍 Predicción Diesel Consumo – Enero 2024:

→ Predicho: 1.31

→ Real: 1317556.51

→ Error Absoluto: 1317555.20

📊 Métricas Diesel Consumo – Enero 2024

MAE: 1317555.20

MSE: 1735951713039.75

RMSE: 1317555.20

```

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/
site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option
adds dropout after all but last recurrent layer, so non-zero dropout ex
pects num_layers greater than 1, but got dropout=0.3 and num_layers=1
warnings.warn(

```

```

In [146... # Usar hiperparámetros óptimos del tuning anterior
best_params_di_imp = df_results_di_imp.iloc[0]
model_di_imp_final = LSTMModel(
    hidden_size=int(best_params_di_imp['hidden_size']),
    dropout=float(best_params_di_imp['dropout'])
).to(device)

```

```

opt_di_imp_final = torch.optim.Adam(
    model_di_imp_final.parameters(),
    lr=float(best_params_di_imp['lr'])
)

# Entrenar modelo final
print("Entrenando modelo final – Diesel Importación...")
history_di_imp_final = train_model(model_di_imp_final, opt_di_imp_final)

# Predicción Enero 2024
model_di_imp_final.eval()
last_window_di_imp = torch.tensor(train_di_imp.values[-12:].reshape(1,

with torch.no_grad():
    pred_di_imp = model_di_imp_final(last_window_di_imp).item()

real_di_imp = test_di_imp.iloc[0]
error_di_imp = abs(pred_di_imp - real_di_imp)

print(f"\n🔴 Predicción Diesel Importación – Enero 2024:")
print(f"→ Predicho: {pred_di_imp:.2f}")
print(f"→ Real:      {real_di_imp:.2f}")
print(f"→ Error Absoluto: {error_di_imp:.2f}")

# Métricas
mae = mean_absolute_error([real_di_imp], [pred_di_imp])
mse = mean_squared_error([real_di_imp], [pred_di_imp])
rmse = mse ** 0.5

print(f"\n📊 Métricas Diesel Importación – Enero 2024")
print(f"MAE:  {mae:.2f}")
print(f"MSE:  {mse:.2f}")
print(f"RMSE: {rmse:.2f}")

```

Entrenando modelo final – Diesel Importación...
Epoch 10, Loss: 884735069388.800049

/Users/gerco/UVG/8th_semester/Data_Science/DC-Lab2/venv/lib/python3.10/site-packages/torch/nn/modules/rnn.py:123: UserWarning: dropout option adds dropout after all but last recurrent layer, so non-zero dropout expects num_layers greater than 1, but got dropout=0.1 and num_layers=1
warnings.warn(

Epoch 20, Loss: 884734414028.800049
Epoch 30, Loss: 884733693132.800049
Epoch 40, Loss: 884733024665.599976
Epoch 50, Loss: 884732382412.800049

📍 Predicción Diesel Importación – Enero 2024:
→ Predicho: 2.13
→ Real: 1348715.64
→ Error Absoluto: 1348713.51

📊 Métricas Diesel Importación – Enero 2024
MAE: 1348713.51
MSE: 1819028137996.83
RMSE: 1348713.51

Comparación final de modelos ARIMA/Prophet

Al incorporar los tres enfoques (ARIMA, Prophet y LSTM), queda claro que la serie de **consumo de gasolina regular** no solo es más “predecible” en términos generales, sino también que el mejor modelo varía según la métrica:

Modelo	Serie	RMSE	MAPE
ARIMA	Importaciones Regular	94 099	7.07 %
	Consumos Regular	49 492	4.39 %
Prophet	Importaciones Regular	107 492	9.88 %
	Consumos Regular	38 262	3.38 %
LSTM	Importaciones Regular	~870 465 ¹	—
	Consumos Regular	~855 349 ¹	—

1. Importaciones Regular

- El **ARIMA** obtuvo aquí un RMSE mucho menor que Prophet (≈ 94 k vs 107 k).
- Sin embargo, el **LSTM**, en el ejercicio puntual, quedó con un error gigantesco (~ 870 k) porque no estaba escalado ni entrenado con la misma granularidad de validación—es decir, para series con grandes rangos absolutos ARIMA/Prophet manejan mejor la escala sin normalizar.

2. Consumos Regular

- El **Prophet** rindió ligeramente mejor que ARIMA (≈ 38 k vs 49 k), ambos con MAPE por debajo del 5 %.
- El **LSTM** también arroja un error absoluto menor en consumos (~ 855 k vs ~ 870 k en importaciones), aunque, de nuevo, la comparación absoluta

con ARIMA/Prophet no es directa por diferencias de escalado.